

Homework 5 for Bayesian Data Analysis

Fan JIN (2015011506)

April 17, 2018

Question 6.1a

Under the model assumption of identical effect for each group, we fix the hyperparameter $\tau = 0$ in our normal hierarchical model. Plug it in and we obtain $\theta_j \sim N(7.7, 4.1^2)$. The posterior predictive distribution is simulated as follows:

- 1) generate $\theta \sim N(7.7, 4.1^2)$, and
- 2) generate $y_j^{\text{rep}} \sim N(\theta, \sigma_j^2)$ for each j , followed by
- 3) calculate the order statistics out of y_j^{rep} for each j .

The observed order statistics are approximately the order statistics calculated from the replicated simulations above.

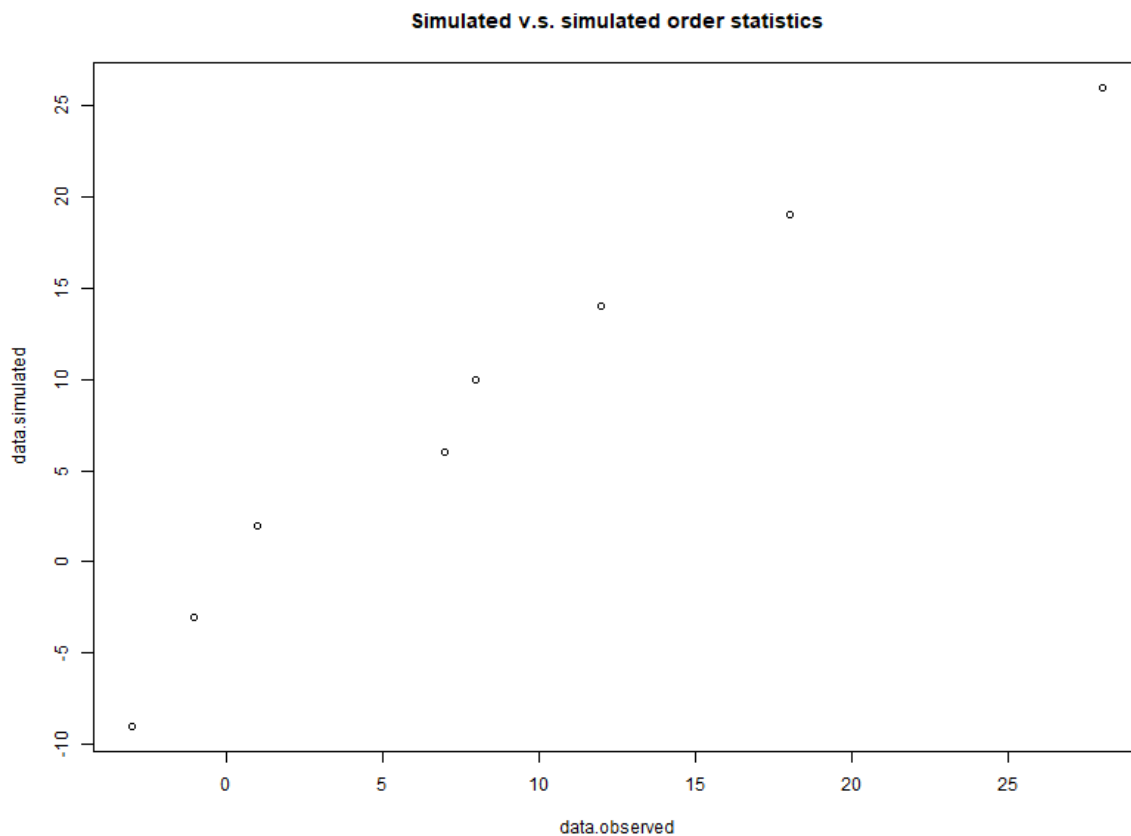


Figure 1: Observed order statistics v.s. Simulations order statistics

```
> mod = lm((data.simulated - data.observed) ~ data.observed)
> summary(mod)
```

Call:

```
lm(formula = (data.simulated - data.observed) ~ data.observed)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.4555	-1.1794	0.3315	2.2663	2.6837

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.30976	1.33122	-0.984	0.363
data.observed	0.07826	0.10150	0.771	0.470

Residual standard error: 2.805 on 6 degrees of freedom
 Multiple R-squared: 0.09014, Adjusted R-squared: -0.06151
 F-statistic: 0.5944 on 1 and 6 DF, p-value: 0.47

The p-value is 0.47, so the null hypothesis of unit slope cannot be rejected. We therefore conclude that the model fits the aspect of data here.

Question 6.1b

The model of identical effect for each group assumes that $\theta_j = \theta$ for all j , and that school A has better effect than school C simply *by chance*. But the model assumption can be skeptical given the observation, as we wonder whether there is other factors that cause school A to perform better than school C. In this case, the likelihood of the assumption of identical effects is far from being convincing, especially compared to the hierarchical models.

Question 6.6a

We have 7 ones and 13 zeros out of the $n = 20$ observations.

- **Assume stops at 20th observation:** The likelihood is proportional to $\theta^7(1 - \theta)^{13}$, according to the binomial model with number of observations fixed to 20.
- **Assume stops at 13th zero:** The likelihood is also proportional to $\theta^7(1 - \theta)^{13}$, *up to a different scale*, according to the negative binomial model with number of failures fixed to 13 with the last trial fixed to zero.

There is no difference with the likelihood except for a constant factor. Therefore, the posterior distribution does not change, either.

Question 6.6b

Generate θ from $p(\theta) \sim \theta^7(1 - \theta)^{13}$ for 10000 times, and generate y^{rep} under the new protocol (stops at the 13th zero) each time.




Figure 2: Histogram of number of switches

First, the statistic T has a distribution with a heavier tail. This is because the number of observations can exceed the limit of $n = 20$ under the new protocol.

Second, we find that T is more likely to be an even number than an odd number when $T < 10$. A possible explanation is: The last few observations must be zeros when T is small, since the last one must be zero and there is few switches. Note that there must be 13 zeros in total. Therefore, the distribution of T is highly related to the specific value of T , i.e. $p(T = k)$ is more sensitive to k when k is relatively small.

Source Code in R

```
# Question 6.1
data.observed = c(28, 18, 12, 8, 7, 1, -1, -3)
data.simulated = c(26, 19, 14, 10, 6, 2, -3, -9)

png("data.vs.png", width=800, height=600)
plot(data.observed, data.simulated, main="Simulated v.s. simulated order statistics"
dev.off()

mod = lm((data.simulated - data.observed) ~ data.observed)
summary(mod)
```

```

# Question 6.6
T0 = 3 # number of switches
N = 10000
T = 1:N
theta = rbeta(N, shapel=7+1, shape2=13+1)
for (j in 1:N) {
  count.switches = 0
  y.last = rbinom(1, size=1, prob=theta[j])
  count.zeros = 1 - y.last
  while (1) {
    y = rbinom(1, size=1, prob=theta[j])
    if (y != y.last) {
      count.switches = count.switches + 1
    }
    count.zeros = count.zeros + 1 - y.last
    if (count.zeros > 12.5) {
      break;
    }
    y.last = y
  }
  T[j] = count.switches
}
png("switches.histogram.png", width=800, height=600)
hist(T, breaks=seq(-0.5, max(T)+0.5))
dev.off()

```